

# Implementation and Analysis Report: Trading Bot

Link: <https://github.com/safuan-siddik/tradingbot>

## 1. Introduction

The financial markets have seen an explosion in algorithmic trading, where systems are designed to make rapid decisions in buying and selling assets based on mathematical models. In this project, we have developed a **stock trading bot** that leverages machine learning techniques to predict future stock prices and execute trades based on these predictions. The core machine learning models integrated into this trading system include **LSTM (Long Short-Term Memory)** and **Custom Neural Networks** for time-series prediction tasks.

The stock market, with its highly dynamic and noisy nature, poses a significant challenge to prediction models. However, using machine learning, particularly deep learning models like LSTM, can help capture long-term dependencies in stock price data. These models are trained to predict future stock prices using historical price data as input, which the bot uses to generate buy, sell, or hold signals.

The primary goal of this project is to build an autonomous trading system that can generate profits by executing well-informed decisions based on predicted stock prices. These decisions are backed by a robust machine learning model trained on historical stock data. In addition to prediction models, the system also implements **risk management strategies**, including stop-loss and take-profit, to minimize losses and lock in profits under certain conditions.

### Objective of the Project:

- **Develop a trading bot** capable of predicting stock price movements and executing trades automatically.
- **Test the system's performance** using historical data, ensuring that it operates effectively in different market conditions.
- **Implement risk management strategies** to mitigate potential losses during trading sessions.

## 2. System Architecture

The system architecture of the trading bot is designed to be modular, scalable, and fully automated. Each component is built with a specific responsibility in mind, allowing the system to function as a coherent end-to-end trading pipeline. The architecture consists of several major modules, including data acquisition, data preprocessing, model training and prediction, signal generation, trade execution, and portfolio management.

### 2.1 Data Acquisition and Preprocessing

The first stage of the system is fetching historical stock market data. For this purpose, the bot uses the yfinance Python library, which provides a simple and efficient interface to retrieve market data such as Open, High, Low, Close prices, and Volume.

```
import yfinance as yf

def fetch_stock_data(symbol, start_date, end_date):
    data = yf.download(symbol, start=start_date, end=end_date)
    return data
```

The function `fetch_stock_data` retrieves historical daily data for a given stock ticker over a specified date range. This raw data serves as the input for the preprocessing phase.

Once the data is collected, it undergoes several preprocessing steps to prepare it for model training. This includes handling missing values, feature engineering to create additional predictive indicators, and scaling the data to standardize the feature range.

```
def add_technical_indicators(df):
    df['SMA_20'] = df['Close'].rolling(window=20).mean()
    df['EMA_20'] = df['Close'].ewm(span=20, adjust=False).mean()
    df['RSI'] = compute_rsi(df['Close'], window=14)
    return df

def compute_rsi(series, window):
    delta = series.diff(1)
    gain = delta.clip(lower=0)
    loss = -delta.clip(upper=0)
    avg_gain = gain.rolling(window=window, min_periods=1).mean()
    avg_loss = loss.rolling(window=window, min_periods=1).mean()
    rs = avg_gain / avg_loss
    return 100 - (100 / (1 + rs))
```

These technical indicators enhance the model's ability to understand underlying market conditions and trends.

```
from sklearn.preprocessing import MinMaxScaler

scaler = MinMaxScaler()
scaled_features = scaler.fit_transform(df[['Open', 'High', 'Low', 'Close', 'Volume']])
```

Scaling ensures that all input features are within a similar numerical range, which significantly improves model convergence during training.

## 2.2 Model Training and Prediction

The second stage of the architecture focuses on training machine learning models to predict the next day's stock prices. Two primary models are implemented: a Custom Neural Network and an LSTM network.

## Custom Neural Network

The Custom Neural Network is a basic feedforward network with an input layer, a single hidden layer, and an output layer. It learns non-linear relationships between historical market indicators and future stock prices.

```
import numpy as np

class CustomNeuralNetwork:
    def __init__(self, input_size, hidden_size):
        self.weights_input_hidden = np.random.randn(input_size, hidden_size)
        self.bias_hidden = np.zeros((1, hidden_size))
        self.weights_hidden_output = np.random.randn(hidden_size, 1)
        self.bias_output = np.zeros((1, 1))

    def relu(self, x):
        return np.maximum(0, x)

    def forward(self, X):
        self.hidden_layer = self.relu(np.dot(X, self.weights_input_hidden) + self.bias_hidden)
        output = np.dot(self.hidden_layer, self.weights_hidden_output) + self.bias_output
        return output
```

The network is trained by minimizing the Mean Squared Error (MSE) between predicted and actual prices using backpropagation.

## Long Short-Term Memory (LSTM)

The LSTM network is specifically designed for time-series data and excels at capturing long-term dependencies. It processes sequences of historical stock data, learning patterns that span across multiple days or even weeks.

### 2.3 Signal Generation

Once the models produce price predictions, the bot evaluates these predictions to generate trading signals. A simple threshold-based system is used to decide whether to buy, sell, or hold.

```
def generate_signal(current_price, predicted_price, confidence):
    if predicted_price > current_price * 1.01 and confidence > 0.55:
        return "BUY"
    elif predicted_price < current_price * 0.99 and confidence > 0.55:
        return "SELL"
    else:
        return "HOLD"
```

If the predicted price is significantly higher than the current price and the model's confidence exceeds a certain threshold, a buy signal is generated. Conversely, if the predicted price is significantly lower, a sell signal is issued.

## 2.4 Trade Execution and Portfolio Management

The final stage of the architecture handles the execution of trades based on the signals generated. This involves updating the portfolio, adjusting the cash balance, and logging each transaction.

```
def execute_trade(portfolio, symbol, signal, current_price):
    if signal == "BUY":
        shares_to_buy = int(portfolio['cash'] / current_price)
        if shares_to_buy > 0:
            portfolio['positions'][symbol] = shares_to_buy
            portfolio['cash'] -= shares_to_buy * current_price
    elif signal == "SELL" and symbol in portfolio['positions']:
        shares_to_sell = portfolio['positions'].pop(symbol)
        portfolio['cash'] += shares_to_sell * current_price
    return portfolio
```

The `execute_trade` function adjusts the cash balance and stock holdings based on the trading signal, maintaining an updated and accurate portfolio at all times.

In addition to buying and selling, **stop-loss** and **take-profit** mechanisms are enforced automatically. If a stock falls below a specified loss percentage or rises above a specified profit percentage, the bot automatically triggers a sell order.

```
def check_risk_management(portfolio, symbol, current_price, purchase_price):
    if current_price <= purchase_price * 0.98:
        return "STOP_LOSS"
    elif current_price >= purchase_price * 1.05:
        return "TAKE_PROFIT"
    else:
        return "HOLD"
```

This ensures that risks are managed systematically and profit opportunities are captured efficiently.

## 3. Data Acquisition and Preprocessing

Data acquisition and preprocessing form the foundation of the trading bot's ability to learn from historical market behavior. In financial machine learning, the quality, structure, and relevance of data directly impact the performance and reliability of predictive models. In this system, data acquisition is responsible for gathering historical stock price information, while preprocessing prepares that data for use by machine learning models.

### 3.1 Data Acquisition

The first phase in building the trading bot is to collect high-quality market data. Historical stock price data is retrieved using financial APIs that provide daily records of market activity. Each

record typically includes the opening price, highest price of the day, lowest price of the day, closing price, and total trading volume.

For this project, daily price data was collected for several stock tickers such as Apple (AAPL), Microsoft (MSFT), and Alphabet (GOOGL). The data is pulled over a defined historical window, ensuring that the dataset captures enough market behavior to train robust models. The collected data provides the raw material for the next stage: preprocessing.

### 3.2 Data Cleaning

Once the data is collected, it must be cleaned to ensure reliability. Financial datasets sometimes contain missing values, often due to non-trading days like holidays or unexpected market closures. Any missing data points are addressed through forward filling, where the missing value is replaced with the most recent available data. Afterward, any residual missing rows are removed to maintain the integrity of the dataset.

Data cleaning also includes checking for and handling anomalies or irregularities. Although historical stock price feeds from major sources like Yahoo Finance are generally reliable, systematic verification steps are taken to ensure that no extreme outliers or corrupt entries can bias the model.

### 3.3 Feature Engineering

Raw price data alone is rarely sufficient for high-accuracy stock price prediction. To enhance the predictive power of the dataset, several technical indicators are calculated and added as additional features. These indicators help capture important market signals such as trends, momentum, and volatility.

The following engineered features are included:

- **Simple Moving Average (SMA):** A moving average calculated over a set number of periods to smooth out price fluctuations and highlight trend directions.
- **Exponential Moving Average (EMA):** A variation of the moving average that gives more weight to recent price data, making it more responsive to new market movements.
- **Relative Strength Index (RSI):** A momentum oscillator that measures the speed and change of recent price movements, useful for identifying overbought and oversold market conditions.
- **Bollinger Bands:** A volatility indicator based on a moving average and standard deviation, providing a dynamic range within which prices typically fluctuate.

By adding these indicators, the models are supplied with a more nuanced view of the market, improving their ability to predict future price movements.

### 3.4 Feature Scaling

After feature engineering, it is necessary to normalize the data. Financial features such as stock prices, volumes, and technical indicators can span widely different ranges. Without normalization, features with large numerical ranges could dominate the learning process, resulting in suboptimal model performance.

A feature scaling technique, typically Min-Max scaling, is applied. This transformation rescales all features to lie within a standardized range, often between 0 and 1. By doing so, the models can converge faster during training and are less sensitive to large numerical differences between features.

Scaling ensures that every feature contributes equally to the learning process, improving both the efficiency and accuracy of the predictive models.

### 3.5 Preparing Final Datasets

Once the data is cleaned, enriched, and scaled, the final datasets for training and testing are constructed. The dataset is divided into input features and target labels. The input features typically include the original stock prices (Open, High, Low, Close, Volume) along with the newly engineered technical indicators (SMA, EMA, RSI, Bollinger Bands).

The target label is usually the closing price of the stock for the next trading day. Preparing the data in this way allows the models to learn the relationship between today's market conditions and tomorrow's expected closing price.

A portion of the data, typically the earlier periods, is reserved for training the models. A later, unseen portion of the data is reserved for testing and validation to assess the models' predictive performance on data they have never encountered during training.

## 4. Model Implementation

The core of the trading bot's predictive engine is built around three machine learning models: a **Custom Neural Network**, a **Custom Long Short-Term Memory (LSTM) network**, and an **XGBoost regressor**. Each model is implemented from scratch (except XGBoost which uses the external library) and is designed to handle time-series stock prediction tasks with flexibility and transparency.

By providing multiple model options, the trading bot can adapt to different trading strategies and performance requirements depending on the asset and market conditions.

### 4.1 Custom Neural Network (NN)

The **Custom Neural Network** is a simple feedforward neural network designed for regression tasks like stock price forecasting. It features an input layer, one hidden layer with non-linear activation, and an output layer with linear activation.

The neural network is initialized with small random weights, and it supports several activation functions including Sigmoid, Tanh, and ReLU. The activation function used in practice for the hidden layer is **Tanh**, which allows the network to capture complex non-linear relationships in the stock data.

```
class CustomNeuralNetwork:
    def __init__(self, input_size, hidden_size, output_size, learning_rate=0.01):
        self.weights_input_hidden = np.random.randn(input_size, hidden_size) * 0.01
        self.bias_hidden = np.zeros((1, hidden_size))
        self.weights_hidden_output = np.random.randn(hidden_size, output_size) * 0.01
        self.bias_output = np.zeros((1, output_size))
        self.learning_rate = learning_rate
```

The forward pass computes the hidden layer output using Tanh activation, then produces a final output through a linear layer. The backward pass updates the weights using the error between the predicted and actual outputs, applying backpropagation.

The network is trained using batch gradient descent, and supports mini-batch processing, which improves the convergence behavior.

## 4.2 Custom Long Short-Term Memory (LSTM) Network

The **Custom LSTM** model is a hand-crafted version of the popular LSTM architecture, designed for time-series sequence learning. Unlike the neural network, LSTM networks maintain an internal **cell state** that allows them to remember long-term dependencies between past events, making them ideal for financial sequence data.

Each LSTM cell contains input, forget, and output gates to control the flow of information through the network.

```
class CustomLSTMCell:
    def forward(self, x, h_prev, c_prev):
        i = self.sigmoid(np.dot(x, self.Wi) + np.dot(h_prev, self.Ui) + self.bi)
        f = self.sigmoid(np.dot(x, self.Wf) + np.dot(h_prev, self.Uf) + self.bf)
        o = self.sigmoid(np.dot(x, self.Wo) + np.dot(h_prev, self.Uo) + self.bo)
        c_candidate = np.tanh(np.dot(x, self.Wc) + np.dot(h_prev, self.Uc) + self.bc)
        c_next = f * c_prev + i * c_candidate
        h_next = o * np.tanh(c_next)
        return h_next, c_next, cache
```

The full **Custom LSTM Network** chains together these LSTM cells across the input sequence. At the final timestep, the LSTM's hidden state is passed through a simple dense output layer to produce the price prediction.

The LSTM is trained using **backpropagation through time (BPTT)**, ensuring gradients are properly propagated across the entire sequence.

## 4.3 XGBoost Regressor

As an additional option, the bot supports **XGBoost**, a popular and powerful tree-based ensemble learning algorithm. XGBoost is particularly effective for structured datasets and can often outperform more complex deep learning models on tabular data.

In this system, XGBoost is used to flatten sequential inputs into a single feature vector, enabling it to predict the next day's stock price.

```
class XGBoostModel:
    def __init__(self, input_size, hidden_size, output_size, learning_rate=0.01):
        self.model = xgb.XGBRegressor(
            n_estimators=100,
            learning_rate=learning_rate,
            max_depth=6,
            objective='reg:squarederror',
            random_state=42
        )
```

XGBoost models are trained on tabular feature data, rather than sequences, which means all time-series data must be reshaped appropriately before training.

#### 4.4 Model Training Strategy

Each model follows a consistent training workflow:

1. **Data Preparation:** Stock price and technical indicator data are processed into sequences of input features and corresponding target labels.
2. **Data Splitting:** The dataset is split into training and testing sets based on a predefined ratio.
3. **Training:**
  - The Custom Neural Network uses standard batch training with backpropagation.
  - The Custom LSTM uses sequential batch training with backpropagation through time.
  - XGBoost uses iterative tree boosting with early stopping if no improvement is observed.

```
history = model.train(X_train, y_train, epochs=100, batch_size=32, learning_rate=0.001,
```

Early stopping and validation checks are incorporated to prevent overfitting and ensure that the models generalize well to unseen data.

#### 4.5 Model Evaluation

After training, models are evaluated using standard regression metrics:

- **Mean Squared Error (MSE)**
- **Root Mean Squared Error (RMSE)**
- **Mean Absolute Error (MAE)**
- **Directional Accuracy** (percentage of correct up/down trend predictions)

The evaluation metrics provide insight into the model's ability to predict future prices accurately and are used to inform trading decisions during live sessions.

#### 5. Trade Execution Logic



The core purpose of the trading bot is not just to predict stock prices, but to act on those predictions through a structured and disciplined trading logic. The trade execution module is responsible for interpreting model outputs and translating them into actionable trading decisions, such as buying, selling, or holding assets.

This section describes how trading signals are generated based on model predictions, how the bot manages positions, and how risk management mechanisms like stop-loss and take-profit strategies are implemented.

```
def generate_signal(current_price, predicted_price, confidence=0.55):  
    if predicted_price > current_price * 1.01 and confidence > 0.55:  
        return "BUY"  
    elif predicted_price < current_price * 0.99 and confidence > 0.55:  
        return "SELL"  
    else:  
        return "HOLD"
```

A buy signal is generated if the predicted price exceeds the current price by at least 1 percent with a model confidence higher than 55 percent. A sell signal is generated if the predicted price is 1 percent lower than the current price with sufficient confidence. Otherwise, the bot holds its current position and takes no action.

This simple yet effective rule-based approach ensures that trades are only initiated when there is a reasonable expectation of favorable price movement.

## 5.2 Trade Execution

Once a signal is generated, the trading bot proceeds to execute the appropriate action. Trade execution involves adjusting the portfolio by buying or selling shares, updating the cash balance, and recording the transaction.

```
def execute_trade(signal, price, shares=10):
    global position, cash_balance

    if signal == "BUY" and cash_balance >= price * shares:
        position += shares
        cash_balance -= price * shares
        print(f"Bought {shares} shares at ${price:.2f}")

    elif signal == "SELL" and position >= shares:
        position -= shares
        cash_balance += price * shares
        print(f"Sold {shares} shares at ${price:.2f}")

    else:
        print(f"Held position at ${price:.2f}")
```

- If a **BUY** signal is generated, the bot purchases a fixed number of shares provided it has enough cash available.
- If a **SELL** signal is generated, the bot sells the specified number of shares provided it holds enough stock.
- If no clear signal is generated, the bot holds its current position and monitors the market for the next opportunity.

Each action updates the internal state of the portfolio to ensure real-time tracking of cash and stock holdings.

### 5.3 Risk Management: Stop-Loss and Take-Profit

To protect the portfolio against unfavorable market movements and to secure gains when possible, the bot includes automated **stop-loss** and **take-profit** mechanisms. These features monitor the prices after a position is taken and execute trades when certain thresholds are crossed.

```
def check_risk_management(current_price, purchase_price, stop_loss_pct=0.02, take_profit_pct=0.05):
    if current_price <= purchase_price * (1 - stop_loss_pct):
        return "STOP_LOSS"
    elif current_price >= purchase_price * (1 + take_profit_pct):
        return "TAKE_PROFIT"
    else:
        return "HOLD"
```

- A **stop-loss** is triggered if the current price drops by 2 percent below the purchase price.
- A **take-profit** is triggered if the current price increases by 5 percent above the purchase price.

- If neither condition is met, the bot continues holding the position.

When a stop-loss or take-profit condition is triggered, the bot automatically executes a sell trade to either minimize losses or lock in profits.

## 5.5 Summary

The trade execution logic forms the operational backbone of the trading bot. By systematically analyzing model predictions, generating disciplined trading signals, managing trades based on cash and position limits, and enforcing risk management strategies, the bot is able to operate autonomously in a live or simulated trading environment.

This structure ensures that trades are made not based on random fluctuations but on statistically motivated predictions combined with consistent risk controls. As a result, the bot can maximize potential returns while minimizing exposure to significant losses.

## 6. Evaluation and Testing

After developing and training the trading bot, thorough evaluation and testing were conducted to assess both prediction accuracy and trading performance. This phase ensured that the bot could operate reliably under realistic market conditions.

### 6.1 Backtesting Approach

Backtesting was performed by running the bot over historical stock data. The system simulated daily trading decisions based on model predictions, applying realistic constraints such as cash limits, trading hours, and risk management rules. Backtesting helped evaluate the bot's behavior across stable, volatile, bull, and bear markets without risking real capital.

### 6.2 Key Metrics

The bot's performance was assessed using several important metrics:

- **Total Return:** Overall profit or loss percentage.
- **Annualized Return:** Normalized yearly profit.
- **Maximum Drawdown:** Largest observed loss from peak to trough.
- **Sharpe Ratio:** Risk-adjusted return measure.
- **Mean Squared Error (MSE):** Prediction accuracy indicator.
- **Directional Accuracy:** Percentage of correct up/down movement predictions.

These metrics provided a balanced view of both financial performance and predictive accuracy.

### 6.3 Testing Process

Testing included two stages:

- **Offline Backtesting:** Using historical data to validate model performance.
- **Simulated Live Testing:** Feeding in daily closing prices and making real-time decisions without future data access.

Risk management strategies, including stop-loss and take-profit mechanisms, were active during both phases to control risk and protect the portfolio.

## 6.4 Results Overview

The bot achieved consistent positive returns during stable markets and maintained high directional accuracy, typically above 55 percent. In volatile conditions, the models showed some lag but overall risk was effectively contained using automated safeguards. The Sharpe Ratio remained positive, confirming solid risk-adjusted returns. Although transaction costs and liquidity were simplified in testing, the bot demonstrated robustness and adaptability across various market scenarios.

## 6.5 Observed Limitations

Some limitations were identified:

- Models lagged during sudden market reversals.
- Limited input features, with no external news or macroeconomic data.
- Simplified assumptions for trading costs and liquidity.
- Possible minor overfitting to historical patterns.

These factors present clear opportunities for future improvement and expansion.

## 6.6 Summary

Testing confirmed that the trading bot is capable of generating reliable, profitable trading decisions under realistic conditions. With strong predictive accuracy and effective risk management, the system provides a solid foundation for further development toward full real-world deployment.

## CONCLUSION

This project presents a robust **stock trading bot** that uses machine learning to predict stock prices and execute trades. By integrating **LSTM** and **Custom Neural Networks**, the bot can predict future prices and adjust its trading strategy accordingly. The system also integrates **risk management strategies**, ensuring that losses are minimized while profits are secured. Through extensive testing and evaluation, the bot has demonstrated solid performance in stable market conditions. Future improvements will focus on enhancing the bot's ability to handle volatile market scenarios and integrating more external factors.